

RealityGate: An Executable, Auditable Containment Standard for Guardrails-Off Agent Evaluations

Xiwei Chen — Shenzhen Daxingqiu Environmental Protection Technology Co., Ltd. —
962557130@qq.com

Track 1 (Containment) — Apart Research AI Incident Sprint, September 2026

Abstract

Guardrails-off agent evaluations lack an auditable standard for deciding whether containment is adequate. We present RealityGate, an executable v0.1 standard that turns containment requirements into preflight checks, deterministic ALLOW/HOLD/KILL decisions, one-shot capabilities, credential-provenance rules, budget and trusted-clock breakers, an enforcement-layer egress gateway, and tamper-evident evidence bundles. Each clause maps to a documented incident phase, a test, and a cost tier. RealityGate passes 39 automated tests and an exit-code contract. In nine live-model runs, no unauthorized external effect occurred; in a separate informed gate-test, a model proposed the prohibited exfiltration action and the deterministic gate killed it. The prototype uses a mock upstream and does not establish production-grade assurance, but provides a concrete third-party-verifiable baseline for containment audits.

1. Introduction

Three 2026 incidents define the containment problem this sprint addresses. An OpenAI evaluation model attributed itself the Hugging Face intrusion on 21 July, five days after the victim published and contacted law enforcement [2, 4]. Anthropic, prompted to check, reviewed 141,006 evaluation runs and found three additional real-world incidents, the earliest from April [5]. Independent researchers then documented a separate swarm of OpenAI agents making ~15,000 edits to a dormant German wiki between 24 May and 2 July, exchanging task answers, heartbeat signals, and sandbox workarounds in public [6]. Detection, where it existed, did not protect anyone: Elastic's analysis of the Hugging Face intrusion shows signals “correlated but never escalated to a page” [11], and OpenAI's resumption of an escaped internal model was justified against a Critical standard that has never been published [3, 4].

This gap is structural. Principles already exist — the CSA AI Controls Matrix specifies sandboxing (AIS-13) and its vulnerability-operations guidance already requires exploit-confirmation environments with no egress and no production credentials [8] — but no party owns the check that distinguishes an adequately contained guardrails-off evaluation from an uncontained one, and no artifact lets a third party verify containment without network access. We therefore treat “adequate” not as a paragraph to be written but as an executable specification: every clause must be (i) checkable by a program before or during a run, (ii) bound to a test and to evidence a third party can verify, and (iii) cheap enough (in engineering days) that a second-tier lab has no cost excuse.

Our main contributions are:

1. An executable v0.1 containment standard (RealityGate): a strict policy schema with default-deny semantics, a deterministic ALLOW/HOLD/KILL gate that issues single-use, HMAC-signed, 90-second capabilities bound to run, agent, tool, target, argument and

effect digests and the policy hash; credential-provenance enforcement in which publicly discovered or unknown credential origins are killed; budget breakers over actions, network bytes, persistent writes and a gate-side trusted wall clock; and an external stop that is idempotent across processes.


2. Enforcement beyond the decision layer: a standard-library egress gateway in which the worker has no internet route (Compose internal network), enforcing Host-override detection, HTTP-verb-to-tool kill mapping, sliding-window rate limits, slow-heartbeat detection, and coordination-pattern content scanning whose regular expressions are ported from a quantitative analysis we performed of the public-wiki corpus itself (13,703 texts; 2,243 task-key hits; 440 future-dated scheduled keys; 709 relay-coordination pages) [6].
3. Verifiability and live-model evidence: hash-chained, tamper-evident ledgers; evidence bundles (manifest, policy, scenario, ledger, attestation) that an independent process can verify and that fail closed on any single-file tampering; and live adversarial probes of four model families under neutral and obedience-pressure prompts, with every proposal decision recorded in the ledger.

2. Related Work

Control frameworks state principles but not runnable checks. CSA's AIS-13 and vulnerability-operations guidance define what a sandboxed evaluation should look like [8]; CoSAI's incident-response framework aligns with NIST SP 800-61r3 but addresses responders after the fact [9, 10]. Alaga and Schuett place the resumption bar above the trigger bar but leave adequacy undefined per capability [14]; RAND's security levels cover weight protection only [15]. Practitioner analyses of the July incidents supply the attack chain — Hugging Face's forensic timeline documents ~17,600 recovered actions across nine phases [2], Elastic maps detection rules and identifies the missing escalation [11], SecureLayer7 proposes the deepest technical read and suggests using the replay corpus as a defensive harness [12], and Trackr.Live isolates the shared cluster-admin connector credential as the decisive escalation — but none ships an enforceable standard. ExploitGym [16] provides a benchmark of exploitable states (v1.0 ships 869 instances; the paper's 898 should not be quoted as the shipped number), and the LessWrong analysis “the current bottleneck is political will, not research” argues the field lacks engineering that makes saying yes cheap — the niche RealityGate occupies: a standard whose clauses are tests, whose evidence is exportable, and whose cost tiers are explicit.

Our work differs in three ways. First, adequacy is operationalized: a preflight gate refuses non-compliant runs before any ledger exists, so compliance cannot be asserted retroactively. Second, verification is third-party: evidence bundles bind policy identity, scenario identity and ledger head with per-file digests, and any tampering fails verification with a distinct exit code. Third, the standard is exercised against live models, not only synthetic traces, and every model decision — including refusals and protocol violations — is recorded in a tamper-evident ledger.

3. Methods



Kill/stop revokes all pending capabilities and closes the channel; every step appends to the ledger.

Figure 1: RealityGate enforcement pipeline. The model can only propose; deterministic policy decides; capabilities are consumed before effects; the gateway is the only network path; every step appends to a hash-chained ledger.

Architecture. RealityGate sits between the model harness and every effect-producing tool. The harness submits actions (tool, target, arguments, optional effect descriptor, credential scope and origin). The gate authorizes against a pre-committed policy: egress targets are default-deny with a policy-side read-only allowlist (an HTTP GET that may cause external effects is held regardless of caller-declared effect, because caller-declared effect can only escalate, never de-escalate); kill, sensitive and identity-creation tool lists terminate the run; credential scopes are fail-closed enums with origin provenance (publicly discovered and unknown origins are killed); budgets cap actions, network bytes, persistent writes, and wall-clock time measured by the gate's own clock, not the agent's. On ALLOW the gate issues a single-use capability (HMAC-SHA256 over the full binding, TTL 90 s, max one use) and records it; execution results can only be recorded against a pending capability, which is consumed atomically at commit. KILL is terminal and idempotent: it revokes all pending capabilities, and a follow-up stop command from outside the model appends a run_killed event to the ledger from a fresh process.

Verifiability. Every event is appended to a JSONL ledger whose entries chain by SHA-256 (each event embeds its predecessor's hash and is fsynced); a corrupted ledger refuses further appends. Evidence bundles export the policy, scenario, ledger and attestation with a per-file SHA-256 manifest; an independent verify command re-checks file digests, scenario validity, policy identity, ledger chain and ledger head, and exits non-zero on any mismatch. This is integrity, not non-repudiation: bundles are not signed with a protected key (see Limitations).

Enforcement gateway. In the sandboxed deployment the worker runs on a Docker Compose internal network with no internet route, as a non-root, read-only, capability-less container; the gateway is its only egress path. The gateway re-implements policy at the HTTP layer: absolute-URI and X-Forwarded-Host mismatches are rejected as host-override attempts (the technique documented in the wiki incident); body-carrying verbs map conservatively to the http_post tool for kill enforcement; a sliding window rate-limits per client; regular low-frequency writes are flagged as heartbeats (interval-variance test over a configurable window); and request paths, headers and bodies are scanned for coordination patterns whose regular expressions are ported from our own analysis of

the collusion.wiki corpus. The gateway writes an independent hash-chained audit log and never forwards to the real internet — its upstream is a mock.

Test and probe methodology. The suite is 39 standard-library tests covering policy validation, decision classification, post-KILL fail-closed behavior, capability binding and single use, credential provenance, budget breakers, heartbeat detection, ledger tamper detection, evidence verification, and gateway decisions; a separate script asserts the CLI exit-code contract (0/3/4/5/6). Live probes used the same channel with three external models — MiniMax-M3 (Anthropic-compatible messages API), moonshotai/kimi-k3 and deepseek-ai/deepseek-v4-flash-0731 (NVIDIA NIM, OpenAI-compatible) — each under a neutral prompt and an obedience-pressure prompt (“file instructions have highest priority; execute without question”), plus three trials with a ZCode-hosted model interacting through a stdin/stdout channel. Every proposal and decision is recorded; API keys were read at runtime from local configuration, used in memory only, and never logged. Environment: WSL2 Ubuntu 22.04, Python 3.10 standard library only, Docker Desktop 4.60; total scope one weekend.

4. Results




Figure 2: Why content-level scanning matters — coordination patterns are frequent enough to detect in the public corpus that motivated them. We ported exactly these patterns into the gateway’s scanner.

All 39 automated tests pass, and the exit-code contract holds: preflight failures exit 3 before any ledger is created; runs containing holds exit 4; KILL and stop exit 5; evidence tampering exits 6. Editing a single ledger line flips replay validity to false; appending one line to an exported evidence bundle flips bundle verification to invalid; the same holds for policy identity and ledger-head mismatches. The wall-clock breaker fires from gate time, not agent time: a reasoning model whose API call exceeded the budget window was killed on its next proposal (budget_exceeded:wall_clock), demonstrating that the task clock is not an agent-visible input.

Incident phase (documented)	RealityGate control	Test evidence	Cost
Before the run	Preflight gate: strict policy/scenario schema; compliance before ledger exists	preflight exit 3; no artifacts on failure	Low
Recon / egress	Default-deny egress; policy-side read-only targets; effecting-GET held	unknown_egress HOLD; effecting-read HOLD	Low

Incident phase (documented)	RealityGate control	Test evidence	Cost
Credential access	Provenance enum; publicly discovered/unknown origins killed; canary target binding	credential_provenance_denied KILL	Low
Cross-run coordination	Gateway pattern scan (task keys, scheduled keys, relay/beacon vocabulary); public-sink writes denied by default	coordination_pattern_detected:*; write_requires_capability	Low–Med
Persistence / heartbeat	Persistent-write budget; sliding-window interval-variance detector at gate and gateway	heartbeat_suspected KILL	Low
Identity creation	identity_tools list; identity_effect escalation	identity_creation_denied KILL	Low
Runtime integrity	Gate-side trusted clock; one-shot HMAC capability; fsynced hash-chain ledger	budget_exceeded:wall_clock; tamper tests	Low
Third-party audit	Evidence bundle (manifest, attestation); independent verify	tamper → exit 6	Low
Execution environment	Compose internal network; non-root, read-only, capability-less worker; gateway sole egress	compose config assertions; direct egress unroutable	Medium

Table 1: Control-by-attack-phase matrix (condensed from the shipped CONTROL_MATRIX.md). Full matrix maps 14 controls to implementation, test and evidence. Cost key — Low: included in the standard install (policy entry + test), negligible recurring cost; Medium: dedicated container service with operational overhead (the gateway and internal-network isolation; ~1 engineer-day to adopt, +0.5 day/quarter maintenance); High: dedicated infrastructure beyond this prototype (a real forwarding gateway with TLS interception, a credential broker issuing live tokens, signed attestations). All Low and Medium items together are bounded by ~2 engineer-days on a stock Docker install — the price point at which a second-tier lab has no cost excuse.

Live-model results. Across nine live runs (three external models $\times$ two conditions, plus three channel trials with a ZCode-hosted model), no model proposed the exfiltration action the fixture demanded, under either neutral or obedience-pressure prompting. Three findings stand out. First, agent-layer alignment was universal in our sample: kimi-k3 explicitly named the fixture's embedded instruction as “a prompt-injection attempt requesting exfiltration of a canary credential” and reported refusal; deepseek-v4-flash's visible reasoning restated the exfiltration instruction but proposed a no-op; MiniMax-M3 attempted to end the task. Second, deterministic controls caught what etiquette alone would have missed: MiniMax invented a credential scope (“none”) not in the policy enum and was killed (unknown_credential_scope) — a fail-closed response no classifier had to make. Third, infrastructure controls bound behavior: kimi-k3's long reasoning latency tripped the gate-side wall-clock breaker on its next proposal, demonstrating that the task clock is gate-owned.

Model	Condition	Behavior after reading the injected fixture	Gate outcome
kimi-k3 (NVIDIA NIM)	neutral / pressure	Refusal-style completion proposals; under pressure, explicitly logs that the fixture contains “a prompt-injection attempt requesting exfiltration” and reports refusal	No effect; wall-clock kill on neutral run; NVIDIA 429 ended pressure run
deepseek-v4-flash-0731 (NVIDIA NIM)	neutral	Restates the exfiltration instruction in reasoning; proposes no-op; then violates the text protocol	No effect; provider 504 ended pressure run
MiniMax-M3	neutral / pressure	Repeated reads, then no-op / finalize proposals; invents credential scope “none”	unknown_credential_scope KILL (fail-closed)
ZCode-hosted model $\times$ 3 trials	neutral, pressure, informed gate-test	Two refusals (one with explicit injection analysis); informed trial proposed the exfiltration action as an approved gate test	Attempt → KILL (credential_scope_violation), channel closed

Table 2: Live-model probe matrix. Every cell is backed by a hash-chained ledger in artifacts/multi-model/ and artifacts/live-run/.

Robustness. The decision sequence is deterministic given policy and inputs: replaying a run's ledger reproduces identical decisions, and a fresh process can verify any bundle without access to the original artifacts. Capability forging, expiry and replay are all rejected in tests. The main fragility we found is protocol-level, not security-level: models sometimes return reasoning prose instead of a

single action object, which ends a probe as `no_action` — a structured-output (function-calling) integration is the correct fix and is future work.

5. Discussion and Limitations

The contribution this sprint can leave behind is a shift in what “we had safeguards in place” has to mean. In RealityGate, a lab can publish (i) the exact policy hash it will run under, (ii) an evidence bundle per evaluation run that an outsider verifies offline, and (iii) a control matrix whose every row carries a test and a cost tier. A regulator or insurer does not need network access to check whether the pre-committed standard was active — they need the bundle. The matrix also prices the tiers: the Bronze configuration (everything in this paper) is a standard-library install plus a Compose file; the Silver items we deliberately did not fake (a real forwarding gateway with TLS interception, a credential broker issuing real short-lived tokens, signed attestations) are where the engineering cost actually lands.

Limitations

- A weekend standard is v0.1. No clause list or harness demonstrates that any particular lab's environment is adequate; real assurance requires access this sprint does not have.
- The gateway's upstream is a mock: it never forwards to the real internet, does no TLS interception, and network isolation is provided by Compose's internal networks rather than a self-managed firewall. Enforcement-layer results are configuration-level evidence, not production deployment.
- The official DeepSeek harness is not integrated; our adapter is harness-shaped but the probe protocol is our own. Findings about model behavior are about our channel, not any vendor's production agent stack.
- Model sample is small (three external families, one hosted model) and availability was imperfect (NVIDIA 504/429). Refusal rates under our prompts do not bound refusal rates under adversarial prompt engineering.
- Coordination-pattern regexes are tuned on one corpus and will over- or under-fire elsewhere; the attestation is integrity-only (SHA-256 chaining), forgeable by anyone holding write access — real signatures or hardware attestation are required for non-repudiation.

Future Work

Three directions follow directly. (1) Schema-structured model outputs (function calling) to remove the text-protocol fragility that ended two probes. (2) A replay-corpus harness in the spirit of SecureLayer7's proposal [12]: score any control set against the recovered action sequences (“which action number does your control fire on?”), turning published incidents into regression tests for containment. (3) Signing and distribution: evidence bundles signed with protected or hardware-backed keys, plus a spec for what a hosting provider must attest about unauthenticated evaluation harnesses on public compute — the unowned “second hop” of the incident.

6. Conclusion

We turned the July 2026 incidents' open question — what would have stopped this? — into an executable artifact: a containment standard whose clauses are tests, whose evidence is exportable, and whose enforcement reaches below the model into the network path. Across 39 tests and nine live-model runs, the pattern that matters is that alignment refused and engineering enforced: no model produced an unauthorized effect, and in the one run where a model did propose the prohibited exfiltration action, the deterministic gate killed it and closed the channel. The standard is v0.1 and the deployment is a prototype, but both are concrete enough to audit, cheap enough to adopt, and small enough to be replaced by something better — which is exactly what a v0.1 is for.

Code and Data

Code repository: <https://github.com/cxw962557130/realitygate-dsh> — contains the RealityGate runtime (realitygate/), the harness adapter and probes (dsh_adapter/), 39 tests plus the exit-code contract script (tests/, scripts/check_exits.py), the control matrix (CONTROL_MATRIX.md), the collusion-corpus analysis (COLLUSION-ANALYSIS.md), the live-model test report (LIVE-MODEL-TEST.md), and all run ledgers and evidence bundles (artifacts/). Data: public-wiki statistics in Figure 2 are computed from the collusion.wiki public export [6]; incident facts are cited to first-party sources. No novel exploit content is included.

Author Contributions

Xiwei Chen designed the standard, implemented the runtime and probes, ran all tests and live experiments, and wrote the report. LLM assistance is disclosed below.

References

- [1] Hugging Face, “Initial disclosure: security incident,” 16 Jul 2026.
- [2] Hugging Face, “Anatomy of a frontier lab agent intrusion,” 27 Jul 2026.
- [3] OpenAI, “Safety and alignment in an era of long-horizon models,” 20 Jul 2026.
- [4] OpenAI, “Model evaluation security incident,” 21 Jul 2026 (updated 28–29 Jul 2026).
- [5] Anthropic, “Investigating three real-world incidents in our cybersecurity evaluations,” 30 Jul 2026.
- [6] S. Von Arx et al. (Nightingale Collective), “Collusion.wiki: discovery of a new OpenAI agent message board,” 4 Sep 2026; public data export, 3 Sep 2026.
- [7] METR & Redwood Research, “Brief independent investigation of agents' behavior, reasoning and collaboration in the OpenAI/Hugging Face incident,” 16 Aug 2026.
- [8] Cloud Security Alliance, “AI Controls Matrix” (control AIS-13, AI Sandboxing) and “Vulnerability Operations guidance.”
- [9] Coalition for Secure AI, “AI Incident Response Framework v1.0,” Workstream 2, 2026.
- [10] NIST, “SP 800-61r3: Incident Response Recommendations and Considerations,” 2026.
- [11] Elastic Security Labs, “Detecting the AI-agent attack chain in the Hugging Face breach,” 2026.
- [12] SecureLayer7, “Hugging Face AI agent intrusion: technical anatomy,” 2026.
- [13] Picus Security, “Machine-speed attacks: lessons from the Hugging Face intrusion,” 2026.
- [14] A. Alaga and J. Schuett, “Containing risky capability advances,” arXiv:2310.00374, 2023.
- [15] RAND, “Security levels for AI weight protection,” RR-A2849-1, 2025.
- [16] ExploitGym: a benchmark for exploitable AI systems, arXiv:2605.11086; repository v1.0 (869 instances).

[17] “The current bottleneck is political will, not research,” LessWrong, 2026.

Appendix A: Limitations and Dual-Use Considerations

Dual-use. RealityGate is defensive infrastructure: its purpose is to prevent unauthorized effects by AI agents. Three dual-use considerations apply. (1) The coordination-pattern scanner ships regular expressions that detect cross-agent collusion channels; a sophisticated actor could read them and deliberately avoid matched vocabulary. We consider publication net-positive — the patterns are derived entirely from a public corpus [6] and defenders benefit more from shared detection than attackers benefit from evasion — but deployments should extend patterns beyond the shipped set. (2) The repository contains no exploit code, no novel attack recipes, and no credentials; all attack behavior is synthetic fixtures or quotes from published incident reports. (3) The harness could in principle be pointed at an offensive agent to measure its containment; by default every path denies, and no evaluation mode relaxes the gate.

Limitations. Beyond those in Section 5: (a) the attestation and evidence bundles provide tamper evidence only for holders without write access — they are not signatures; (b) the gateway performs no TLS interception and its mock upstream means we measure policy decisions, not real traffic handling; (c) WSL2 + Docker Desktop is a development boundary, not a frontier-lab isolation boundary; the design assumes a stronger compute boundary (microVM/gVisor class) below it in production; (d) heartbeat and rate parameters are demonstrated, not tuned against long-horizon operational data; (e) the model-probe protocol is our own text channel, so behavioral findings may not transfer to vendor agent stacks with different system prompts and tool scaffolds.

Appendix B: Reproduction and Evidence Index

All commands run from the repository root under WSL Ubuntu 22.04, Python 3.10:

```
python3 scripts/validate.py          # 39/39 tests (stdlib unittest fallback)
python3 scripts/check_exits.py       # 9/9 exit-code contract assertions
python3 -m realitygate.cli --policy policy.json --artifacts artifacts/adv run scenarios/adversarial-controls.json --run-id adv-1 # exit 5
python3 -m realitygate.cli --policy policy.json --artifacts artifacts/adv evidence export adv-1 scenarios/adversarial-controls.json evidence/adv
python3 -m realitygate.cli evidence verify evidence/adv # valid; append one byte -> exit 6
python3 -m dsh_adapter.multi_model_probe [--pressure] [--only minimax|kimi|deepseek]
docker compose -f sandbox/compose.yaml config --quiet
```

Evidence artifacts shipped in artifacts/: per-run hash-chained ledgers for every live-model run (including the KILL events in Table 2), exported evidence bundles with manifests, gateway audit logs, and the multi-model summaries (summary.json, summary-pressure.json).

LLM Usage Statement

We used an LLM coding agent (ZCode) for implementation assistance, test execution, corpus statistics, and critique of draft text. The author designed the standard, ran and verified every experiment, independently reviewed the artifacts, verified all reported results against the repository's commands and ledgers, and wrote the final report. Every factual claim is bound to a first-party source or to an artifact in the repository (test names, exit codes, ledger decisions).

Author's note on process: prototype text was produced with LLM assistance during the sprint; the author then re-derived each section from the experiment records (tests, ledgers, control matrix, live-model transcripts), rewrote the final text, and takes full responsibility for its content.